



深度學習系統與實現

LAB5

Deployment of Trained Models

Dept. of Computer Science and
Information Engineering

National Chiao Tung University



LAB 5

- ❑ Dataset: Food11
- ❑ Model: Modified Model
- ❑ Accuracy: 91+ on evaluation dataset
- ❑ Realized with TensorRT



Outline

- Tool
- TensorRT workflow
- LAB 5-1
- LAB 5-2
- Report Spec
- Grading




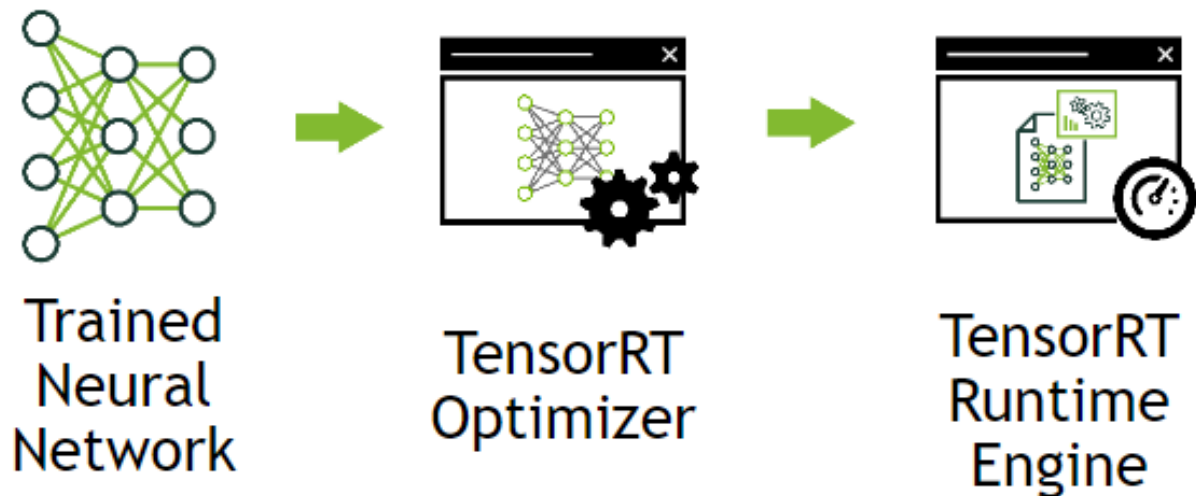
Tool

□ TensorRT

- Official Docs: <https://docs.nvidia.com/deeplearning/sdk/tensorrt-developer-guide/index.html>
- Support two interfaces
 - Python API (you will also need pyCUDA)
 - C++ API
- Using docker image are highly encouraged
 - <https://docs.nvidia.com/deeplearning/sdk/tensorrt-container-release-notes/running.html>
 - Be aware of each release's TensorRT and CUDA version

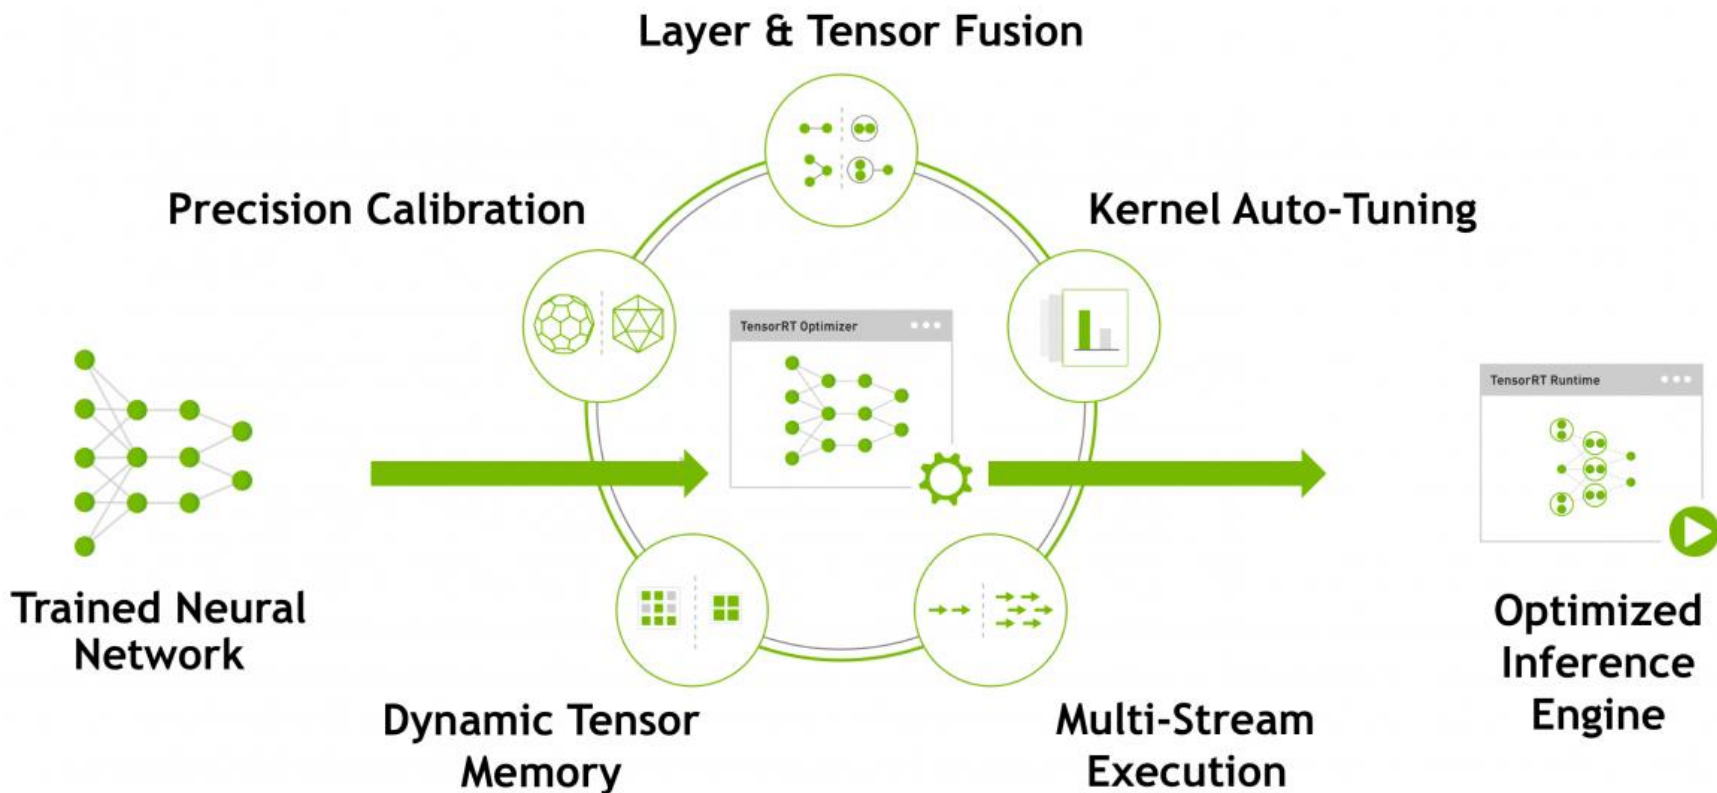


TensorRT Workflow



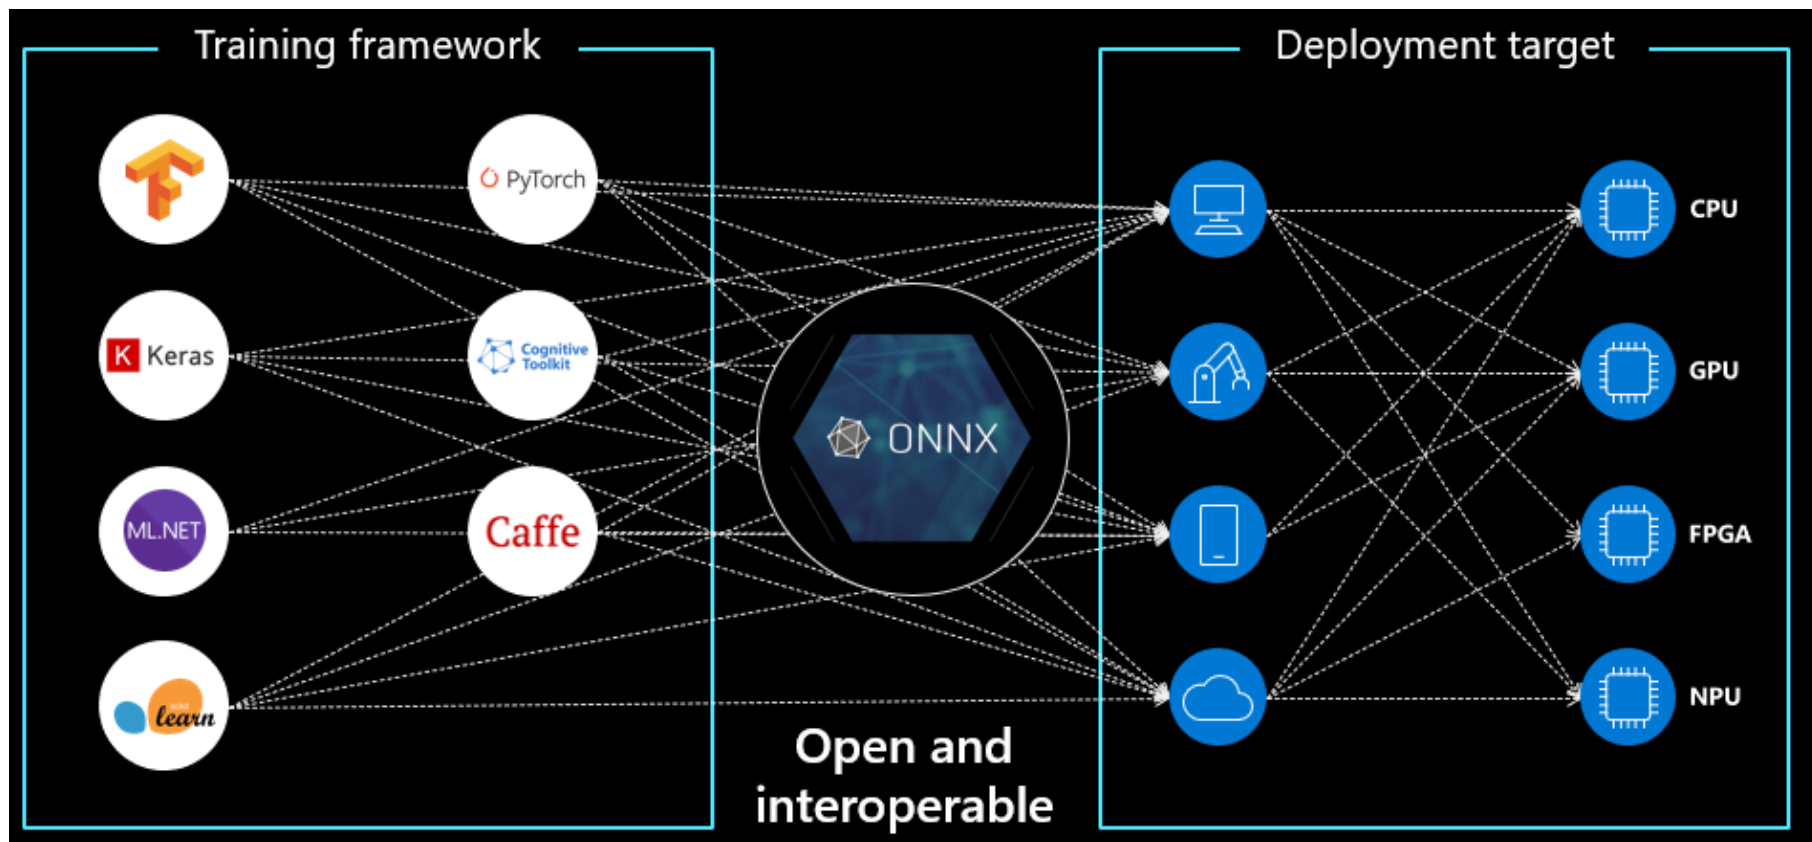


TensorRT Optimizer





Open Neural Network Exchange (ONNX)



Credit: <https://microsoft.github.io/ai-at-edge/docs/onnx/>



TensorRT workflow (ONNX Parser)

- Prepare builder, network definition, and parser
 - We use TensorRT ONNX parser here

```
builder = trt.Builder(TRT_LOGGER)
network = builder.create_network()
parser = trt.OnnxParser(network, TRT_LOGGER)

with open(model_file, 'rb') as model:
    parser.parse(model.read())
```

- Build engine via the network from ONNX parser

```
builder.build_cuda_engine(network)
```
- Create context from engine

```
with engine.create_execution_context() as context:
```




Perform engine on CUDA devices

- Prepare Input/output on CPU/GPU
 - Based on test image size and batch size

- Copy input to GPU

- Perform inference

- Copy output to CPU

```
import tensorrt as trt
```

```
TRT_LOGGER = trt.Logger(trt.Logger.WARNING)  
ONNX_MODEL = "mnist.onnx"
```

```
def build_engine():
```

```
    with trt.Builder(TRT_LOGGER) as builder, builder.create_network() as network, \  
        trt.OnnxParser(network, TRT_LOGGER) as parser:  
        # Configure the builder here.  
        builder.max_workspace_size = 2**30  
        # In this example, we use the ONNX parser, but this should be replaced  
        # according to your needs. This step might instead use the Caffe/UFF parser,  
        # or even the Network API to build a TensorRT Network manually .  
        with open(ONNX_MODEL, 'rb') as model:  
            parser.parse(model.read())  
        # Build and return the engine. Note that the builder,  
        # network and parser are destroyed when this function returns.  
        return builder.build_cuda_engine(network)
```

```
def do_inference():
```

```
    with build_engine() as engine, engine.create_execution_context() as context:  
        # Allocate buffers and create a CUDA stream before inference.  
        # This should only be done once.  
        pass  
        # Preprocess input (if required), then copy to the GPU, do inference,  
        # and copy the output back to the host.  
        pass
```



Important Parameters

- WorkspaceSize
 - GPU **temporary memory** which the engine can use at **execution time**.
- MaxBatchSize
 - The maximum batch size which can be used at execution time, and also the **batch size** for which the engine will be **optimized**.



LAB 5-1(50%)

- Train target model to target accuracy
 - Use any technique to train the model to target accuracy
 - Target Accuracy = 91% ↑
- Serialize to ONNX
 - Serialize the model to ONNX
 - The serialize detail please reference the link below
 - <https://pytorch.org/docs/stable/onnx.html#functions>



LAB 5-2(50%)

- Inference on TensorRT
 - Use the ONNX from 5-1
- Batch Size Adjustment
 - Set the batch size to [1, 2, 4, 8, 16, 32, 64]
 - Show the latency and FPS in plot graph



Report

- Environment Setup
 - GPU 、 CUDA Version 、 TensorRT Version
 - Baseline Setup(batch size, pre-processing etc...)
- Baseline
 - Accuracy and Inference speed
- Result
 - TensorRT accuracy on evaluation dataset
 - Batch size vs Speedups
- Conclusion and Insights
- Anything you want to say



Grading

- LAB 5-1 (50%)
 - Train model to target accuracy (30%)
 - Serialize to ONNX (20%)
- LAB 5-2 (50%)
 - Inference on TensorRT (40%)
 - Batch size adjustment (10%)
- Bonus (10%)
 - Further Optimization
 - Multiple stream on TensorRT
 - Mixed precision inference
 - et cetera...
- Submission to E3:
 - Source code + report
 - zip format (ex: dllab_lab5_{group id}.zip)
- Deadline : 2020/05/25 23:55

Total:
110



Reference

- PyTorch ONNX API
 - <https://pytorch.org/docs/stable/onnx.html>
- TensorRT:
 - <https://docs.nvidia.com/deeplearning/sdk/tensorrt-developer-guide/index.html>
- TensorRT API Documentation
 - https://docs.nvidia.com/deeplearning/sdk/tensorrt-api/c_api/index.html